

Deep Physics-aware Inference of Cloth Deformation for Monocular Human Performance Capture

Yue Li^{1,2,3}, Marc Habermann^{1,2}, Bernhard Thomaszewski³, Stelian Coros³, Thabo Beeler⁴, and Christian Theobalt^{1,2}

¹Max Planck Institute for Informatics, ²Saarland Informatics Campus, ³ETH Zürich, ⁴Google Inc.

Abstract

Recent monocular human performance capture approaches have shown compelling dense tracking results of the full body from a single RGB camera. However, existing methods either do not estimate clothing at all or model cloth deformation with simple geometric priors instead of taking into account the underlying physical principles. This leads to noticeable artifacts in their reconstructions, such as baked-in wrinkles, implausible deformations that seemingly defy gravity, and intersections between cloth and body. To address these problems, we propose a person-specific, learning-based method that integrates a finite element-based simulation layer into the training process to provide for the first time physics supervision in the context of weakly-supervised deep monocular human performance capture. We show how integrating physics into the training process improves the learned cloth deformations, allows modeling clothing as a separate piece of geometry, and largely reduces cloth-body intersections. Relying only on weak 2D multi-view supervision during training, our approach leads to a significant improvement over current state-of-the-art methods and is thus a clear step towards realistic monocular capture of the entire deforming surface of a clothed human.

1. Introduction

Human performance capture plays a critical role in various computer graphics and vision applications such as virtual try-on, avatar re-targeting, movies as well as video games. With rapid progress in display and capture technology, expectations on the quality of geometric reconstruction and tracking are constantly increasing. Here, not only the geometric details are of major importance but also that the deformed and posed reconstructions follow the physical behavior of real objects which includes realistic wrinkle patterns as well as coherent interaction of body and clothing. While professional content production studios can rely on involved multi-camera setups to capture high-fidelity hu-

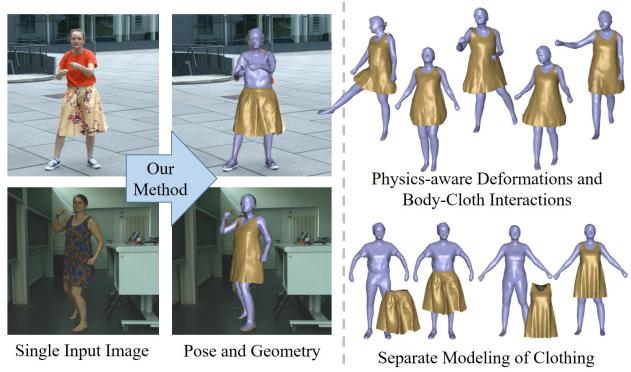


Figure 1: Our method estimates body pose as well as physically plausible surface deformation from a single image. Importantly, body and clothing are represented as separate meshes which allows accurate modeling of body-cloth interactions.

man performances, there is an ever growing desire to democratize performance capture for everyday applications, such as virtual try-on, by using much simpler and cheaper capture devices.

Hence, research has shifted from expensive and complex multi-view capture setups [52, 6, 10, 7, 8, 50, 14, 17, 38, 44, 72, 73, 77, 50] to depth cameras [61, 22, 46, 28, 21, 36, 80, 16, 82, 83, 81, 76] over the past decade. Unfortunately, the latter are sensitive to bright sunlight and thus not suited for outdoor use-cases. Due to this and the advances in deep learning, the most recent research has shifted its attention onto *single RGB camera* setups, offering the most flexible and low cost setup. Previous monocular methods have made a substantial progress in recovering the 3D unclothed body [31, 49, 32], hand pose [74, 43, 88] as well as facial identity and expression [34, 66, 67]. First attempts have also been made to jointly track the above body parts [48, 78, 30]. However, so far only few methods [24, 23, 79] have managed to coherently track the dense surface deformations with clothing included from monocular views, which is essential for a lot of applications and for an immersive experience. These person-specific methods densely deform and pose a geometry to match the body pose as well as the cloth-

ing deformation in the input image while assuming an initial template of the person is given. Recent learning-based monocular methods [24] only assume image-based supervision, making it challenging for them to densely supervise deformations. This manifests in simplified model assumptions like a single geometry for body and clothing as well as the usage of simple geometric priors which do not obey physics principles. As a result, they cannot account for body-cloth interactions, deformations do not follow physical rules, and they have notable artifacts such as baked-in wrinkles from the initial scan.

To this end, we propose a learning based approach for capturing body pose as well as the physically plausible clothing deformation from a single RGB image (see Fig. 1). Our method is comprised of two networks dedicated to regress body pose in terms of joint angles and surface deformations in form of embedded deformation. Importantly, during training we only assume weak supervision in form of multi-view imagery, i.e. 2D skeletal joint detection and foreground masks. As discussed before, this type of supervision can hardly ensure physically plausible results. Thus, at the core of our approach, we propose an efficient simulation layer that for the first time allows physically plausible self supervision *during* the training in such a weakly supervised setting. We achieve this by integrating a finite element based simulator into a learning architecture which takes intermediate predictions of cloth and body positions and velocities to perform a forward simulation. The simulation result is then used to supervise the deformations of our deformation network during training. As cloth-body collisions are explicitly handled in the proposed layer, we can accurately model clothing as a separate piece of geometry in contrast to previous monocular methods. In summary, our contributions are:

- A monocular human performance capture approach, which outputs body pose and physically plausible cloth deformations for dressed subjects.
- A simulation network layer that allows on-the-fly simulation supervision *during* training, which also enables separate modeling of cloth and body geometry.

In contrast to prior work, our method reconstructs physically more accurate deformations without baked in wrinkles and with correct body-cloth collision handling. Our quantitative evaluations indicate that incorporating physical simulation during training provides significant improvements over state-of-the-art methods.

2. Related Work

As our goal is recovering a dense surface of the human, we focus on previous works that achieve this by using parametric body models or template meshes, and works that

treat body and clothing as separate mesh layers. We do not discuss the large body of work on 2D [11, 12, 60, 75] and 3D skeletal pose estimation [42, 41, 25, 54, 53, 87, 65, 69, 55] as they are not concerned with the problem of surface reconstruction.

Reconstruction of Parametric Body Models. The works [86, 29, 56, 19, 4, 37, 31, 70] that fall into this category use parametric body models such as SMPL [39]. Some works fit the model parameters to sparse 2D and 3D joint predictions [4] or to regressed vertex positions [35] by minimizing corresponding energies. Others [31] directly regress these parameters from images. A set of recent works [48, 78] extended body models to account for changing hand pose and facial expression which was then used to jointly capture hands, face and body. While motion and shape of the naked body can be reconstructed, clothing was not considered.

Unified Reconstruction. One stream of previous work treats body and clothing as a single geometry. Volumetric representations [85, 71] use an occupancy grid to represent the body, meaning that resolution is limited by the grid. Implicit methods [57, 27, 58] methods overcome this limitation by treating the subject’s surface as an implicit function. However, both approaches require post-processing to recover an explicit surface representation, e.g., through the marching cubes algorithm. The resulting concerns in terms of temporal consistency make these approaches all but impractical for applications such as texture replacement or motion retargeting. Closely related to our work are template-based methods [23, 24, 79, 13, 15] which track a template based on image observations. Using a mesh with fixed topology as reference, surface correspondence over time is directly given. With input data originating from images only, however, they must replace missing information with simplifying assumptions and geometric priors that do not reflect real physical behavior. Geometric priors alone cannot ensure physically plausible result, leading to static wrinkles present in template that remain visible across all poses, and deformations that do not respect gravity. Most importantly, all these methods treat clothing and body as a single piece of geometry so that they cannot account for dynamic body-cloth interactions. To address these limitations, we propose a simulation layer that encourages cloth deformations to not only satisfy image constraints, but to also exhibit physically plausible behavior.

Cloth as a Separate Part. In contrast to the above methods, there is also a line of work that reconstructs body and clothing as separate geometries. Bhatnagar et al. [3] recover static geometry for clothing and body from a set of RGB images. ClothCap [51] uses multi-view capture to

produce a clothed human body that can be used for re-targeting. Stoll et al. [63] recover cloth material parameters from multi-view video sequences to reproduce the observed garment deformation. Similar to our work, but requiring a depth camera, is SimulCap [84], which performs quasi-static physics simulation with depth matching constraints to reconstruct the clothing layer. While all of the above methods rely either on multi-view setups or depth cameras, our approach requires only a single RGB camera.

As a potential alternative to simulation, geometric detail such as wrinkles can be added in a data-driven, pose-dependent manner [18, 59, 47, 20]. Different from these geometry-driven methods, we integrate physics-based simulation into our training framework thus encouraging physical plausibility with only a single image as input.

3. Method

Our template-based method leverages a deep neural architecture, which takes a single background-segmented person image as input and regresses posed and deformed surface meshes for body and clothing which match the performance in the input image (see Fig. 2). Before training, a 3D template of the person with separate cloth and body geometry as well as a multi-view recording of the subject performing various motions has to be acquired (Sec. 3.1). The technical core of our architecture is formed by two prediction networks, *PoseNet* and *SimNet*, that are trained to regress body pose and cloth deformation, respectively (Sec. 3.3). *PoseNet*, introduced by Habermann et al. [24], regresses skeleton joint angles and the root rotation from the input image using multi-view 2D joint detections as weak supervision. The proposed *SimNet* predicts the surface deformation of the cloth template by regressing embedded graph parameters from the same input image. In addition to multi-view image data, *SimNet* leverages our cloth simulation layer as supervision, which encourages physically plausible deformations (Sec. 3.2).

3.1. Data Processing

Template Acquisition. Similar to [24], we acquire a single scan for body and clothing (e.g. using photogrammetric scanning). We perform surface registration against a parametric body mesh model using the method of Bhatnagar et al. [1, 2] to obtain an estimate for body parts occluded by clothing, such as the legs under a skirt, which are merged with the visible body parts from the scan to form a complete body mesh. We also mark the arms of the body mesh as inactive when resolving collisions in our simulation. A separate cloth mesh is created manually from the scan, a task that could also be automated [63, 51]. Skeleton parameters and skinning weights, required for posing the two meshes, are determined automatically as described in [24]. Two separate embedded graphs [64, 62] for body and cloth-

ing are computed by subsampling the original body and cloth meshes. For more details, we refer to the supplemental document.

Video Capture. We capture the subject to be tracked in a multi-view green screen studio with calibrated and synchronized cameras. We ask the person to perform various tasks such as walking and dancing to best sample the space of possible poses. Next, we run OpenPose [12, 11, 60, 75] on all frames and views to obtain multi-view 2D joint predictions. Color keying is used to segment the foreground from the greenscreen background and compute distance transform images D_c from the foreground masks [5].

3.2. Cloth Simulation Layer

Our simulation layer uses the publicly-available cloth simulation framework ARCSim [45] as its basis, but we make several adjustments that we describe below.

Material Model and Parameter Selection. ARCSim uses a data-driven material model defined by a total of 39 parameters. While parameter values for several real-world fabrics are provided, we found that none of them were ideally suited for the materials that we use in our examples and manually adjusting parameters to obtain better approximation proved very difficult. For this reason, we resorted to a simpler, isotropic material model (see, e.g., [68]) that is defined through three parameters: Young's modulus and Poisson's ratio for in-plane behavior, and a single bending stiffness coefficient. We determine parameter values through best-guess initialization and a few iterations of simulation-based tuning to better approximate the qualitative behavior observed in the input video sequences. Although we found this manual approach to be sufficient for our examples, this task could be further automated [63].

Time Integration. During training, the initial state and velocities that are fed into the simulation layer can exhibit large deformations that, when using ARCSim's default integration method, can lead to instabilities. To improve stability, we resort to an optimization-based formulation of fully implicit Euler (see, e.g., [40]) combined with adaptive regularization and a back-tracking line search. Finally, we make several code adaptations to enable batch operations for efficient training and integrate the simulation engine in a customized TensorFlow¹ layer. We refer to this layer as the simulation function S later, which takes cloth and body vertices as input and returns the cloth positions for the next time step.

¹<https://www.tensorflow.org/>

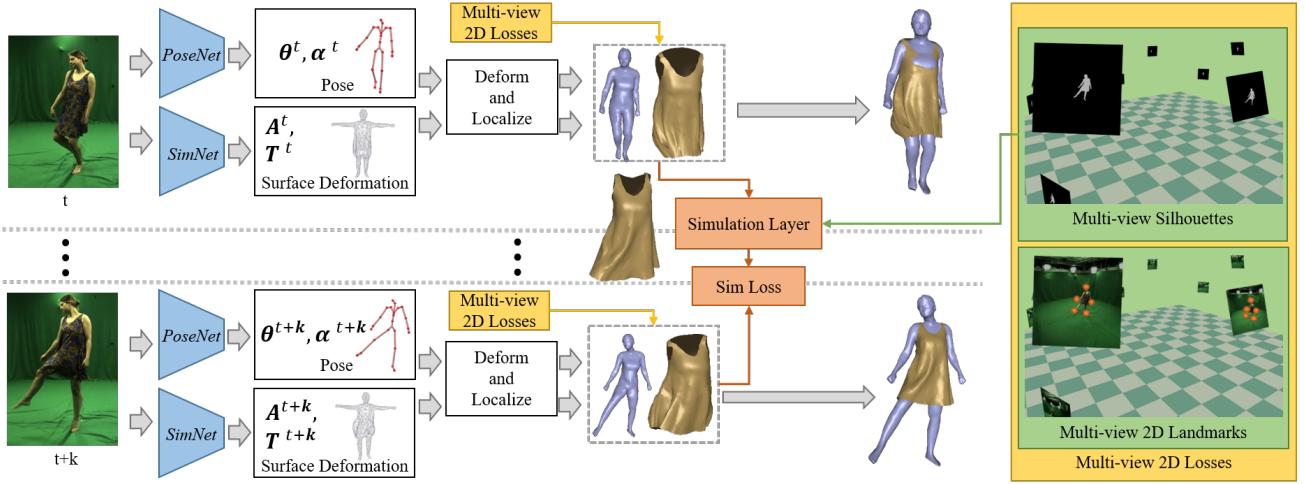


Figure 2: Overview. Our method takes a single image as input and two separate networks, *PoseNet* and *SimNet*, regress the skeletal pose as well as embedded deformation parameters for the clothing. Combining the outputs of the two networks allows posing and deforming the body and clothing geometry. During training, we use multi-view image losses for *PoseNet*, whereas *SimNet* is additionally supervised by our proposed simulation loss to encourage physically plausible deformations. To evaluate the simulation loss, we run on-the-fly cloth simulation on small windows of subsequent frames from the training sequence and penalize the difference between regressed deformations and simulation outputs.

Silhouette Constraint. Although our simulation model captures the characteristic behavior of clothing, it is still an approximation and deviations from the input images must be expected due to external forces such as air drag, viscous damping, and friction that are not modeled. To better track the real-world behavior, we therefore introduce an additional multi-view silhouette constraint, integrated into the simulation layer. Specifically, this constraint ensures that the vertices $\tilde{\mathbf{V}}_{\text{cloth}}^t$ of the simulated cloth geometry matches the image silhouettes from all camera views for frame t . Therefore, we construct a 3D ray going through the camera origin and the silhouette pixel p and search for the boundary vertex $\tilde{\mathbf{V}}_{\text{cloth},p}^t$ that minimizes the distance to this ray. The closest point on the ray is used as 3D point correspondence for the boundary vertex, included as soft constraint into the simulation.

$$E_{\text{cons}} = \sum_p \|\tilde{\mathbf{V}}_{\text{cloth},p}^t - \mathbf{V}_{\text{ray},p}^t\|^2. \quad (1)$$

Next, we introduce our pose regressor and the surface deformation network which leverages our simulation layer as an additional form of supervision.

3.3. Pose and Deformation Regression

We separate the task of regressing the full surface deformation into predicting pose and surface deformation independently. Therefore, our method consists of two ResNet50 based CNNs [26], *PoseNet* and *SimNet*, which regress skeleton pose and embedded deformation parameters from a segmented input image, respectively.

3.3.1 Pose Regression and Deformation Model

To pose and deform template vertices as well as sparse body markers, the deformation layer of Habermann et al. [24] denoted as

$$\mathbf{V}_{\text{loc}}, \mathbf{K}_{\text{loc}} = f(\theta, \alpha, \mathbf{A}, \mathbf{T}) \quad (2)$$

is used, which is a combination of dual quaternion skinning [33] and embedded deformation [64, 62]. It takes the pose in terms of skeleton joint angles $\theta \in \mathbb{R}^3$ and camera-relative root joint rotation $\alpha \in \mathbb{R}^3$ as well as the embedded graph node rotation $\mathbf{A} \in \mathbb{R}^{K \times 3}$ and translation $\mathbf{T} \in \mathbb{R}^{K \times 3}$ where each row encodes rotations in terms of Euler angles and translation vectors for each of the K nodes. The output is the posed and deformed vertices \mathbf{V}_{loc} and markers \mathbf{K}_{loc} in camera and root relative space where the i th row contains the updated position of vertex i . The body pose parameters θ and α are regressed using *PoseNet* [24].

3.3.2 Physics-aware Deformation Regression

To not only pose the template but also account for surface deformation, we propose a dedicated network *SimNet* which predicts the translation vectors \mathbf{T} and rotation angles \mathbf{A} of the embedded graph (EG) from the segmented input image. *SimNet* is supervised using a combination of both image-based and physics-based metrics, which ensure that deformations match image-based observations while minimizing violations of physical equilibrium conditions. In the remainder of this chapter we assume *PoseNet* is fixed and provides the posed and deformed vertices \mathbf{V} and markers \mathbf{K} in global space. As \mathbf{V} and \mathbf{K} are a function of the *Sim-*

Net output (\mathbf{T} , \mathbf{A}), we can supervise SimNet by supervising \mathbf{V} and \mathbf{K} .

Warm Start. To jump start our training including the simulation layer, we first pre-train *SimNet* without running simulation but use a geometric regularizer (ARAP [62]). This adds robustness to the training since geometric regularizers are more stable than simulation and significantly reduces overall training time. Once the network predicts reasonable shapes, we add the simulation loss to supervise the physical deformation. The loss is defined as

$$L_{\text{warm}} = L_{\text{sil}} + L_{\text{lm}} + L_{\text{reg}} + L_{\text{att}} \quad (3)$$

which comprises multi-view losses as well as geometric priors. The individual loss terms are defined as follows.

Multi-view Image Losses. First, a multi-view 2D landmark loss

$$L_{\text{lm}} = \sum_c \sum_m \|\Pi_c(\mathbf{K}_m) - \mathbf{p}_{c,m}\|^2 \quad (4)$$

is used, which ensures that projected landmark matches the 2D detection $\mathbf{p}_{c,m}$ for all views c and landmarks m . Here, Π_c denotes the projection function of camera c . To densely supervise the surface, we also introduce a silhouette loss

$$L_{\text{sil}} = \sum_c \sum_{b \in \mathcal{B}_c} \rho_{c,b} \|\Pi_c(\mathbf{V}_b) - \mathbf{D}_c\|^2, \quad (5)$$

which ensures that the set of mesh boundary vertices \mathcal{B}_c matches the zero contour line in the distance transform image \mathbf{D}_c for all camera views. $\rho_{c,b}$ is a weighting term ensuring that silhouettes are only matched if the normal of the surface aligns with the gradient of the distance transform [23].

Regularization Loss. To regularize deformations and to avoid drifting of the surface, we use the as-rigid-as-possible prior [62] which ensures that local embedded deformations are smooth. We further adopt the original formulation by using rigidity weights [24] allowing us to model material dependent deformation behaviours, e.g. the skirt can deform more freely than the skin.

Attachment Loss. Note that our entire mesh \mathbf{V} can be split into body and garment meshes, denoted as $\mathbf{V}_{\text{cloth}}$ and \mathbf{V}_{body} in the remainder of this section. To ensure a coherent movement of the two, an attachment loss

$$L_{\text{att}} = \sum_{i \in \mathcal{A}} \|\mathbf{V}_{\text{cloth},i} - \sum_{j=0}^2 \gamma_{i,j} \mathcal{C}(\mathbf{V}_{\text{cloth},i}, \mathbf{V}_{\text{body}})_j\|^2 \quad (6)$$

is included, which ensures that the cloth is attached to the body at some anchor positions, e.g. the waistband of a skirt has to be attached to the hip of the body mesh. Here, \mathcal{A} are the selected vertices on the garment that act as anchor points, \mathcal{C} is a function that takes the cloth vertex id i and returns the 3 vertices of its closest triangle on the undeformed body mesh, and γ_{ij} are precomputed barycentric weights computed from the closest point on this triangle and its three vertices.

Physics-aware Training. Although our previous training allows the surface to already match image evidence, it can neither account for collision of body and clothing nor ensure physically plausible cloth deformations. To this end, we introduce an additional simulation-based loss function, whose goal it is to explicitly supervise collision behaviour as well as physical plausibility. Substituting the ARAP loss by the simulation loss, our final loss is defined as

$$L = L_{\text{sil}} + L_{\text{lm}} + L_{\text{sim}} + L_{\text{att}}. \quad (7)$$

As our previously introduced simulation layer \mathcal{S} is directly integrated into a learning framework, we can perform on-the-fly simulation during training. As training samples we use short frame sequences of length \mathcal{F} starting at random frames t' . In the following, we refer to a specific frame in this window using the superscript \cdot^t , where $t \in \{t', \dots, t' + \mathcal{F}\}$. Our physics loss

$$L_{\text{sim}} = \sum_i \sum_{t=t'+1}^{t'+\mathcal{F}} \|\mathbf{V}_{\text{cloth},i}^t - \tilde{\mathbf{V}}_{\text{cloth},i}^t\|^2 \quad (8)$$

now ensures that deformed cloth vertices $\mathbf{V}_{\text{cloth},i}^t$ match the simulation result for all frames within the frame window except the first one. Here, $\tilde{\mathbf{V}}_{\text{cloth},i}^t$ are the simulated cloth vertex positions which are defined as

$$\tilde{\mathbf{V}}_{\text{cloth}}^t = \begin{cases} \mathbf{V}_{\text{cloth}}^t, & t = t' \\ \mathcal{S}(\mathbf{V}_{\text{cloth}}^{t-1}, \mathbf{V}_{\text{cloth}}^t, \mathbf{V}_{\text{body}}^{t-1}, \mathbf{V}_{\text{body}}^t), & t = t' + 1 \\ \mathcal{S}(\tilde{\mathbf{V}}_{\text{cloth}}^{t-2}, \tilde{\mathbf{V}}_{\text{cloth}}^{t-1}, \mathbf{V}_{\text{body}}^{t-1}, \mathbf{V}_{\text{body}}^t), & t > t' + 1 \end{cases}$$

where \mathcal{S} takes current and previous body and cloth positions to approximate garment and body velocities. Note that for $t = t'$ no velocities are available, thus we set $\tilde{\mathbf{V}}_{\text{cloth}}^t = \mathbf{V}_{\text{cloth}}^t$. Similarly for $t = t' + 1$, we have to approximate cloth velocity using $\mathbf{V}_{\text{cloth}}^{t-1}$, $\mathbf{V}_{\text{cloth}}^t$. For the remaining frames in the window, the velocities can be acquired from previous simulation steps using forward differences (FD). We can now evaluate the simulation loss on these $\mathcal{F}-1$ frames (excluding the first frame of the window) and thus supervise our *SimNet* deformation predictions. We opted to not backpropagate gradients through the simulation inputs with respect to the EG parameters as this would

make it harder to guarantee convergence during training. Even though the first frame in a training sample sequence does not receive supervision, that frame is supervised by our multi-view supervision, such that in practice all frames are supervised. In addition to the shape of the cloth, the networks also generate the body shape. In the next section, we demonstrate that incorporating physics during training provides stronger and more physically plausible supervision leading to significantly less collisions, and physically implausible forces.

4. Results

We evaluate our approach on various outdoor and indoor environment settings with three subject-cloth combinations under a wide range of motions (see Fig. 3). To bridge the domain gap between training data recorded in the capture studio and in-the-wild testing sequences, e.g. different light conditions and cameras, we apply a domain adaptation step. Here, PoseNet and SimNet are refined for 300 iterations on the testing sequence leveraging the losses introduced before but *using only a single camera*. For in-the-wild captures with varying and dynamic backgrounds, we segment the input images using OSVOS [9]. While the result is almost collision free thanks to *SimNet*, minute intersections can still remain, which is why we run a final collision resolution step (see also supplemental video). Following [24], we smooth the final results temporally using a Gaussian filter of size 5 frames. To visualize our result overlayed onto the input image, we leverage the single view global alignment of Habermann et al. [24] that computes the global translation in closed form.

Dataset. Our training dataset contains 3 green screen studio capture sequences ($S4$, $F1$, $F2$) with actors performing a large range of motions. For testing, we recorded an additional multi-view green screen sequence to evaluate our reconstruction on reference views and multiple in-the-wild captures using a single camera with a resolution of 1920×1080 for every subject. Sequence $S4$ is from [24], and we additionally acquired two training sequences and templates, $F1$ and $F2$, with 18 cameras at a resolution of 1285×940 , where each sequence contains around 20,000 frames. We will release the dataset for future research.

Qualitative Results. In Fig. 3, we test our method on various in-the-wild environments while the subjects perform a wide range of motions. Our method does not only provide accurate image overlays and plausible 3D body and cloth geometries but our reconstruction also shows physics-aware cloth deformations and plausible body-cloth interactions. In Fig. 4, we show that our SimNet predicts different physically plausible wrinkle patterns related to the charac-

ter motion. This is due to our separate modeling of body and cloth geometry and the fact that body-cloth interactions are taken into account by our simulation supervision during training. We further visualize the underlying body geometry without clothing. Note that also the occluded parts are predicted accurately.

4.1. Comparisons

We compare our approach to the state-of-the art template-based monocular human performance capture methods [23, 24]. LiveCap [23] determines poses using inverse kinematics with predicted 2D and 3D joint positions as targets and computes surface deformations such as to optimize a single-view image loss. DeepCap [24] uses weak supervision from multi-view images during training to predict pose and embedded deformation parameters from a single segmented image.

4.1.1 Qualitative Comparisons

In Fig. 5, we compare our method with state-of-the-art template-based methods [23, 24]. Unlike our approach, both LiveCap and DeepCap only use geometric priors on the deformations during optimization and training, respectively. Consequently, the resulting cloth deformation contains static wrinkles from the initial template (see top left corners) that persist across all poses. By using simulation supervision and separate modeling of cloth and body geometries, the wrinkles generated by our method are less constrained by the template and, consequently, exhibit more variety and better physical plausibility.

4.1.2 Quantitative Comparisons

We evaluate our results using the green screen testing sequence of $S4$ for all metrics below. For a fair comparison, we use the same cloth-body geometry, obtained through manual clean-up of the input scans, for all approaches.

Out-of-balance Force Evaluation. In Tab. 1, we list the magnitude of the out-of-balance forces, which are defined as the difference between inertial forces and the sum over internal, external, and collision forces. This physical measure indicates to what extent the results deviate from Newton’s second law of motion, and vanishes for physically-correct motion. The acceleration of the body and the garment for a given frame is determined using a centered difference approximation based on network predictions for three consecutive frames. To reduce the global translation error irrespective of our network predictions, we apply the ground truth global translation for all methods as described by Habermann et al. [24]. Our method performs not only better on average compared to other approaches but also



Figure 3: Reconstructions obtained with our method for various in-the-wild environments and challenging motion combinations. Our results show good overlay quality throughout, attesting to the pose and clothing estimation accuracy of our method. Furthermore, diverse and physically plausible cloth deformations are observed for a wider range of poses.

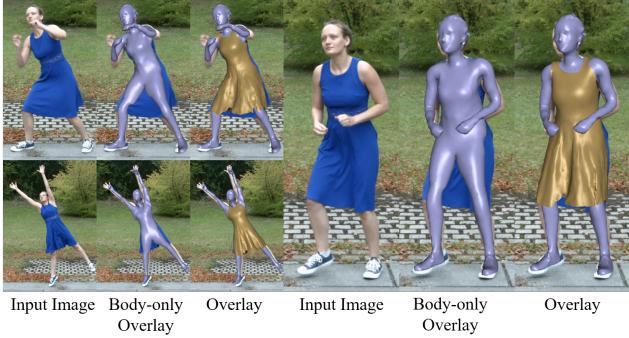


Figure 4: Physically plausible deformation. Here we show our reconstructed body geometry as well as our final model with body geometry and clothing. Despite minor distortions in occluded regions, our reconstructed body geometry matches the image evidence. Thanks to the separate template geometry modeling, our cloth reconstruction is able to reproduce the folding and unfolding behaviour of the dress which is driven by the underlying body motions.

significantly reduces the peak value. LiveCap [23] performs significantly worse due to the inherent ambiguity of the single-image setting combined with the inability of geometric priors to capture physical behavior. As a result, the cloth geometry returned by their method exhibits large distortions, in particular in regions occluded from view, resulting in large internal forces. DeepCap [24] leverages neural

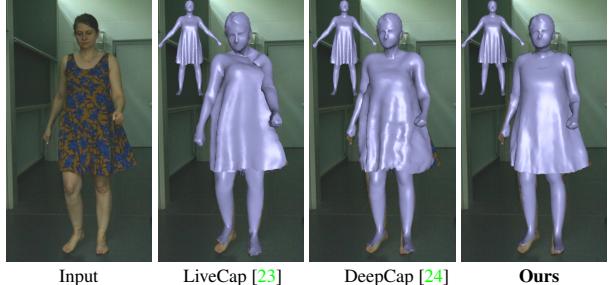


Figure 5: Comparison with state-of-the-art template-based monocular methods [23, 24]. Using simulation supervision during training, our method produces physically more realistic results without baked-in wrinkles from the initial template mesh (shown in the top left corners).

network models trained with multi-view supervision. Despite substantial improvements compared to LiveCap, our physics-aware method leads to another 50% decrease in error.

Garment-body Intersection Distance. In Tab. 2, we further compare the garment-body intersection distance to LiveCap [23] and DeepCap [24]. Their simplified one-piece templates sidestep this issue, however, when we evaluate with our more accurate mesh model, e.g. separate geometry for clothing and body, collisions severely affect the reconstruction quality. We show that our network predictions significantly reduce cloth penetration.

Methods	Avg	Max
LiveCap[23]	79.85	1140
DeepCap[24]	2.119	29.84
Ours	1.017	7.063

Table 1: Out-of-balance force evaluation. We compare our method to LiveCap [23] and DeepCap [24] with respect to out-of-balance force magnitude. It can be seen that our physics-aware method outperforms state-of-the art geometry-based methods.

Average Penetration Depth (cm)	
Methods	Distance
LiveCap[23]	25.58
DeepCap[24]	23.83
Ours	4.165

Table 2: Penetration Depth. We compute average penetration depths for cloth-body intersections across a 10,000 frame testing sequence. Not taking collisions into account, both LiveCap [23] and DeepCap [24] produce severe penetrations. Our method handles collisions during training, which leads to substantially reduced penetration depths.

IoU Percentage. To measure reconstruction quality from different camera views, we compare our results with previous methods using the intersection over union (IoU) metric (see Tab. 3). The IoU metric indicates the overlapping percentage of the camera projection images of our reconstruction and the foreground segmentation of input images (ground truth). To be consistent with [24], the evaluation is performed for every 100th frame of the testing sequence from S4. We apply the same ground truth global translation and temporal filter as in DeepCap [24]. Comparing to body-only reconstruction methods, our method achieves significantly better performance. It should be noted that, compared to the other approaches, the cloth geometry in our method is more constrained due to physics. For example, the strap of a dress cannot detach from the body to match the image silhouettes. Nonetheless, we achieve comparable IoU accuracy while maintaining better physical plausibility.

Methods	AMVIoU (%)	RVIoU(%)	SVIoU(%)
HMR[31]	65.10	64.66	70.84
LiveCap[23]	59.96	59.02	72.16
DeepCap[24]	82.53	82.22	86.66
Ours	80.83	80.53	84.83

Table 3: IoU percentage comparison. Average multi view (*AMVIoU*), reference view (*RVIoU*) and single view (*SVIoU*) values correspond to IoU evaluation on all views, all views expect input view, and input view, respectively. Our reconstruction provides comparable accuracy with state-of-the-art methods while delivering more physically plausible results.

4.2. Ablation Study

Simulation during Training. Here we verify that, with simulation supervision in the training process, physically unrealistic cloth deformations and other artefacts resulting from merely image-based supervision can be reduced. As can be seen in Fig. 6, the strap of the dress remains on the body, and the bottom of the dress does not distort to match the silhouette.

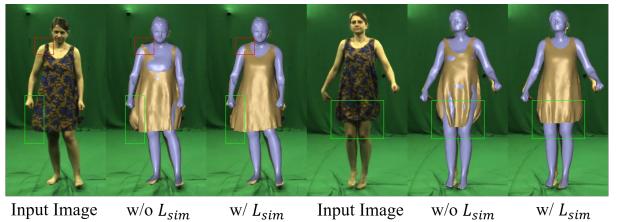


Figure 6: Simulation *during* training. We evaluate our deformation network with and without simulation loss on a testing sequence. It can be seen that penetrations and deformation artefacts are largely reduced when using simulation.

5. Conclusion

We propose a physics-aware deep learning-based method for monocular human performance capture. With physics-based simulation running on-the-fly as a network layer, we enforce physics plausibility to compensate the shorthand of using only multi-view images. We show more visually pleasing results and much improved physics-metrics over state-of-the-art methods.

Limitations & Future Work. Simulation is the computational bottleneck of our training process. However, using GPU-accelerated simulation, we believe that training time can be substantially reduced. On a related note, the simulation layer is not differentiable in our current implementation. Nevertheless, a fully-differentiable physics solver would improve data efficiency and it would open the door to automatic material parameter estimation from video input. Our method is able to faithfully track body pose and cloth deformations for dynamic input motion, but it cannot produce dynamic effects from a single input image—an inherently ill-posed problem. To further improve physical fidelity and reconstruction quality, we would like to extend our method to regress dynamically-consistent cloth motion by leveraging deep temporal architectures, which take short videos as input instead of single frames.

References

- [1] Thiendo Alldieck, Marcus Magnor, Bharat Lal Bhatnagar, Christian Theobalt, and Gerard Pons-Moll. Learning to reconstruct people in clothing from a single RGB camera. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, jun 2019.
- [2] Bharat Lal Bhatnagar, Cristian Sminchisescu, Christian Theobalt, and Gerard Pons-Moll. Combining implicit function learning and parametric models for 3d human reconstruction. In *European Conference on Computer Vision (ECCV)*. Springer, August 2020.
- [3] Bharat Lal Bhatnagar, Garvita Tiwari, Christian Theobalt, and Gerard Pons-Moll. Multi-garment net: Learning to dress 3d people from images. In *IEEE International Conference on Computer Vision (ICCV)*. IEEE, oct 2019.
- [4] Federica Bogo, Angjoo Kanazawa, Christoph Lassner, Peter Gehler, Javier Romero, and Michael J. Black. Keep it SMPL: Automatic estimation of 3D human pose and shape from a single image. In *European Conference on Computer Vision (ECCV)*, 2016.
- [5] Gunilla Borgefors. Distance transformations in digital images. *Computer Vision, Graphics, and Image Processing*, 34(3):344 – 371, 1986.
- [6] Matthieu Bray, Pushmeet Kohli, and Philip HS Torr. Posecut: Simultaneous segmentation and 3d pose estimation of humans using dynamic graph-cuts. In *European conference on computer vision*, pages 642–655. Springer, 2006.
- [7] Thomas Brox, Bodo Rosenhahn, Daniel Cremers, and Hans-Peter Seidel. High accuracy optical flow serves 3-d pose tracking: exploiting contour and flow based constraints. In *European Conference on Computer Vision*, pages 98–111. Springer, 2006.
- [8] Thomas Brox, Bodo Rosenhahn, Juergen Gall, and Daniel Cremers. Combined region and motion-based 3d tracking of rigid and articulated objects. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 32(3):402–415, 2010.
- [9] Sergi Caelles, Kevis-Kokitsi Maninis, Jordi Pont-Tuset, Laura Leal-Taixé, Daniel Cremers, and Luc Van Gool. One-shot video object segmentation. In *Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [10] Cedric Cagniart, Edmond Boyer, and Slobodan Ilic. Free-form mesh tracking: a patch-based approach. In *Computer Vision and Pattern Recognition (CVPR), 2010 IEEE Conference on*, pages 1339–1346. IEEE, 2010.
- [11] Zhe Cao, Gines Hidalgo, Tomas Simon, Shih-En Wei, and Yaser Sheikh. OpenPose: realtime multi-person 2D pose estimation using Part Affinity Fields. In *arXiv preprint arXiv:1812.08008*, 2018.
- [12] Zhe Cao, Tomas Simon, Shih-En Wei, and Yaser Sheikh. Realtime multi-person 2d pose estimation using part affinity fields. In *CVPR*, 2017.
- [13] Joel Carranza, Christian Theobalt, Marcus A Magnor, and Hans-Peter Seidel. Free-viewpoint video of human actors. *ACM transactions on graphics (TOG)*, 22(3):569–577, 2003.
- [14] Edilson De Aguiar, Carsten Stoll, Christian Theobalt, Naveed Ahmed, Hans-Peter Seidel, and Sebastian Thrun. Performance capture from sparse multi-view video. In *ACM Transactions on Graphics (TOG)*, volume 27, page 98. ACM, 2008.
- [15] Edilson De Aguiar, Carsten Stoll, Christian Theobalt, Naveed Ahmed, Hans-Peter Seidel, and Sebastian Thrun. Performance capture from sparse multi-view video. In *ACM SIGGRAPH 2008 papers*, pages 1–10. 2008.
- [16] Mingsong Dou, Sameh Khamis, Yury Degtyarev, Philip Davidson, Sean Ryan Fanello, Adarsh Kowdle, Sergio Orts Escalano, Christoph Rhemann, David Kim, Jonathan Taylor, et al. Fusion4d: Real-time performance capture of challenging scenes. *ACM Transactions on Graphics (TOG)*, 35(4):114, 2016.
- [17] Juergen Gall, Carsten Stoll, Edilson De Aguiar, Christian Theobalt, Bodo Rosenhahn, and Hans-Peter Seidel. Motion capture using joint skeleton tracking and surface estimation. In *Computer Vision and Pattern Recognition, 2009. CVPR 2009. IEEE Conference on*, pages 1746–1753. IEEE, 2009.
- [18] Peng Guan, Loretta Reiss, David A Hirshberg, Alexander Weiss, and Michael J Black. Drape: Dressing any person. *ACM Transactions on Graphics (TOG)*, 31(4):1–10, 2012.
- [19] Peng Guan, Alexander Weiss, Alexandru O Bălan, and Michael J Black. Estimating human shape and pose from a single image. In *ICCV*, pages 1381–1388, 2009.
- [20] Erhan Gundogdu, Victor Constantin, Amrollah Seifoddini, Minh Dang, Mathieu Salzmann, and Pascal Fua. Garnet: A two-stream network for fast and accurate 3d cloth draping. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 8739–8748, 2019.
- [21] Kaiwen Guo, Jonathan Taylor, Sean Fanello, Andrea Tagliasacchi, Mingsong Dou, Philip Davidson, Adarsh Kowdle, and Shahram Izadi. Twinfusion: High framerate non-rigid fusion through fast correspondence tracking. 09 2018.
- [22] Kaiwen Guo, Feng Xu, Tao Yu, Xiaoyang Liu, Qionghai Dai, and Yebin Liu. Real-time geometry, albedo, and motion reconstruction using a single rgb-d camera. *ACM Transactions on Graphics (TOG)*, 36(3):32, 2017.
- [23] Marc Habermann, Weipeng Xu, Michael Zollhoefer, Gerard Pons-Moll, and Christian Theobalt. Livecap: Real-time human performance capture from monocular video. *ACM Transactions on Graphics (TOG)*, 38(2):1–17, 2019.
- [24] Marc Habermann, Weipeng Xu, Michael Zollhoefer, Gerard Pons-Moll, and Christian Theobalt. Deepcap: Monocular human performance capture using weak supervision. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5052–5063, 2020.
- [25] Ikhsanul Habibie, Weipeng Xu, Dushyant Mehta, Gerard Pons-Moll, and Christian Theobalt. In the wild human pose estimation using explicit 2d features and intermediate 3d representations, 2019.
- [26] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- [27] Zeng Huang, Yuanlu Xu, Christoph Lassner, Hao Li, and Tony Tung. Arch: Animatable reconstruction of clothed humans. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3093–3102, 2020.

- [28] Matthias Innmann, Michael Zollhöfer, Matthias Nießner, Christian Theobalt, and Marc Stamminger. VolumeDeform: Real-time Volumetric Non-rigid Reconstruction. October 2016.
- [29] Arjun Jain, Thorsten Thormählen, Hans-Peter Seidel, and Christian Theobalt. MovieReshape: Tracking and reshaping of humans in videos. *ACM Transactions on Graphics*, 29(5), 2010.
- [30] Hanbyul Joo, Tomas Simon, and Yaser Sheikh. Total capture: A 3d deformation model for tracking faces, hands, and bodies. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018.
- [31] Angjoo Kanazawa, Michael J. Black, David W. Jacobs, and Jitendra Malik. End-to-end recovery of human shape and pose. In *Computer Vision and Pattern Recognition (CVPR)*, 2018.
- [32] Angjoo Kanazawa, Jason Y Zhang, Panna Felsen, and Jitendra Malik. Learning 3d human dynamics from video. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 5614–5623, 2019.
- [33] Ladislav Kavan, Steven Collins, Jiří Žára, and Carol O’Sullivan. Skinning with dual quaternions. In *Proceedings of the 2007 Symposium on Interactive 3D Graphics and Games*, I3D ’07, 2007.
- [34] Hyeongwoo Kim, Pablo Garrido, Ayush Tewari, Weipeng Xu, Justus Thies, Matthias Nießner, Patrick Pérez, Christian Richardt, Michael Zollöfer, and Christian Theobalt. Deep video portraits. *ACM Transactions on Graphics (TOG)*, 37(4):163, 2018.
- [35] Nikos Kolotouros, Georgios Pavlakos, and Kostas Daniilidis. Convolutional mesh regression for single-image human shape reconstruction, 2019.
- [36] Adarsh Kowdle, Christoph Rhemann, Sean Fanello, Andrea Tagliasacchi, Jonathan Taylor, Philip Davidson, Mingsong Dou, Kaiwen Guo, Cem Keskin, Sameh Khamis, David Kim, Danhang Tang, Vladimir Tankovich, Julien Valentin, and Shahram Izadi. The need 4 speed in real-time dense visual tracking. In *SIGGRAPH Asia 2018 Technical Papers*, SIGGRAPH Asia ’18, pages 220:1–220:14, New York, NY, USA, 2018. ACM.
- [37] Christoph Lassner, Javier Romero, Martin Kiefel, Federica Bogo, Michael J. Black, and Peter V. Gehler. Unite the people: Closing the loop between 3d and 2d human representations. In *Proc. CVPR*, 2017.
- [38] Yebin Liu, Carsten Stoll, Juergen Gall, Hans-Peter Seidel, and Christian Theobalt. Markerless motion capture of interacting characters using multi-view image segmentation. In *Computer Vision and Pattern Recognition (CVPR), 2011 IEEE Conference on*, pages 1249–1256. IEEE, 2011.
- [39] Matthew Loper, Naureen Mahmood, Javier Romero, Gerard Pons-Moll, and Michael J Black. Smpl: A skinned multi-person linear model. *ACM transactions on graphics (TOG)*, 34(6):1–16, 2015.
- [40] Sebastian Martin, Bernhard Thomaszewski, Eitan Grinspun, and Markus Gross. Example-based elastic materials. *ACM Trans. Graph.*, 30(4), July 2011.
- [41] Dushyant Mehta, Oleksandr Sotnychenko, Franziska Mueller, Weipeng Xu, Mohamed Elgharib, Pascal Fua, Hans-Peter Seidel, Helge Rhodin, Gerard Pons-Moll, and Christian Theobalt. XNect: Real-time multi-person 3D motion capture with a single RGB camera. volume 39, 2020.
- [42] Dushyant Mehta, Srinath Sridhar, Oleksandr Sotnychenko, Helge Rhodin, Mohammad Shafiei, Hans-Peter Seidel, Weipeng Xu, Dan Casas, and Christian Theobalt. Vnect: Real-time 3d human pose estimation with a single rgb camera. volume 36, 2017.
- [43] Franziska Mueller, Florian Bernard, Oleksandr Sotnychenko, Dushyant Mehta, Srinath Sridhar, Dan Casas, and Christian Theobalt. Ganerated hands for real-time 3d hand tracking from monocular rgb. In *Proceedings of Computer Vision and Pattern Recognition (CVPR)*, June 2018.
- [44] Armin Mustafa, Hansung Kim, Jean-Yves Guillemaut, and Adrian Hilton. General dynamic scene reconstruction from multiple view video. In *ICCV*, 2015.
- [45] Rahul Narain, Armin Samii, and James F O’brien. Adaptive anisotropic remeshing for cloth simulation. *ACM transactions on graphics (TOG)*, 31(6):1–10, 2012.
- [46] Richard A. Newcombe, Dieter Fox, and Steven M. Seitz. Dynamicfusion: Reconstruction and tracking of non-rigid scenes in real-time. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2015.
- [47] Chaitanya Patel, Zhouyingcheng Liao, and Gerard Pons-Moll. Tailornet: Predicting clothing in 3d as a function of human pose, shape and garment style. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7365–7375, 2020.
- [48] Georgios Pavlakos, Vasileios Choutas, Nima Ghorbani, Timo Bolkart, Ahmed A. A. Osman, Dimitrios Tzionas, and Michael J. Black. Expressive body capture: 3d hands, face, and body from a single image, 2019.
- [49] Georgios Pavlakos, Luyang Zhu, Xiaowei Zhou, and Kostas Daniilidis. Learning to estimate 3d human pose and shape from a single color image, 2018.
- [50] Gerard Pons-Moll, Sergi Pujades, Sonny Hu, and Michael Black. ClothCap: Seamless 4D clothing capture and retargeting. *ACM Transactions on Graphics, (Proc. SIGGRAPH)*, 36(4), 2017.
- [51] Gerard Pons-Moll, Sergi Pujades, Sonny Hu, and Michael J Black. Clothcap: Seamless 4d clothing capture and retargeting. *ACM Transactions on Graphics (TOG)*, 36(4):1–15, 2017.
- [52] Gerard Pons-Moll, Javier Romero, Naureen Mahmood, and Michael J Black. Dyna: a model of dynamic human shape in motion. *ACM Transactions on Graphics (TOG)*, 34(4):120, 2015.
- [53] Alin-Ionut Popa, Mihai Zanfir, and Cristian Sminchisescu. Deep multitask architecture for integrated 2d and 3d human sensing. *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [54] Helge Rhodin, Jörg Spörri, Isinsu Katircioglu, Victor Constantin, Frédéric Meyer, Erich Müller, Mathieu Salzmann, and Pascal Fua. Learning monocular 3d human pose estimation from multi-view images, 2018.
- [55] Gregory Rogez, Philippe Weinzaepfel, and Cordelia Schmid. Lcr-net: Localization-classification-regression for human

- pose. In *CVPR 2017-IEEE Conference on Computer Vision & Pattern Recognition*, 2017.
- [56] Lorenz Rogge, Felix Klose, Michael Stengel, Martin Eisemann, and Marcus Magnor. Garment replacement in monocular video sequences. *ACM Transactions on Graphics (TOG)*, 34(1):6, 2014.
- [57] Shunsuke Saito, Zeng Huang, Ryota Natsume, Shigeo Morigima, Angjoo Kanazawa, and Hao Li. Pifu: Pixel-aligned implicit function for high-resolution clothed human digitization. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 2304–2314, 2019.
- [58] Shunsuke Saito, Tomas Simon, Jason Saragih, and Hanbyul Joo. Pifuhd: Multi-level pixel-aligned implicit function for high-resolution 3d human digitization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 84–93, 2020.
- [59] Igor Santesteban, Miguel A Otaduy, and Dan Casas. Learning-based animation of clothing for virtual try-on. In *Computer Graphics Forum*, volume 38, pages 355–366. Wiley Online Library, 2019.
- [60] Tomas Simon, Hanbyul Joo, Iain Matthews, and Yaser Sheikh. Hand keypoint detection in single images using multiview bootstrapping. In *CVPR*, 2017.
- [61] Mirosława Slavcheva, Maximilian Baust, Daniel Cremers, and Slobodan Ilic. Killingfusion: Non-rigid 3d reconstruction without correspondences. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, volume 3, page 7, 2017.
- [62] Olga Sorkine and Marc Alexa. As-rigid-as-possible surface modeling. In *Symposium on Geometry processing*, volume 4, pages 109–116, 2007.
- [63] Carsten Stoll, Juergen Gall, Edilson De Aguiar, Sebastian Thrun, and Christian Theobalt. Video-based reconstruction of animatable human characters. *ACM Transactions on Graphics (TOG)*, 29(6):1–10, 2010.
- [64] Robert W Sumner, Johannes Schmid, and Mark Pauly. Embedded deformation for shape manipulation. In *ACM SIGGRAPH 2007 papers*, pages 80–es. 2007.
- [65] Xiao Sun, Jiaxiang Shang, Shuang Liang, and Yichen Wei. Compositional human pose regression. *ICCV*, 2017.
- [66] Ayush Tewari, Michael Zollhöfer, Pablo Garrido, Florian Bernard, Hyeongwoo Kim, Patrick Pérez, and Christian Theobalt. Self-supervised multi-level face model learning for monocular reconstruction at over 250 hz. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018.
- [67] Ayush Tewari, Michael Zollöfer, Hyeongwoo Kim, Pablo Garrido, Florian Bernard, Patrick Perez, and Theobalt Christian. MoFA: Model-based Deep Convolutional Face Autoencoder for Unsupervised Monocular Reconstruction. In *The IEEE International Conference on Computer Vision (ICCV)*, 2017.
- [68] Bernhard Thomaszewski, Simon Pabst, and Wolfgang Straßer. Asynchronous cloth simulation. In *Computer Graphics International*, volume 2, page 2, 2008.
- [69] Denis Tome, Chris Russell, and Lourdes Agapito. Lifting from the deep: Convolutional 3d pose estimation from a single image. *IEEE Conf. on Computer Vision and Pattern Recognition. Proceedings*, 2017.
- [70] Gü̈l Varol, Duygu Ceylan, Bryan Russell, Jimei Yang, Ersin Yumer, Ivan Laptev, and Cordelia Schmid. BodyNet: Volumetric inference of 3D human body shapes. In *ECCV*, 2018.
- [71] Gü̈l Varol, Duygu Ceylan, Bryan Russell, Jimei Yang, Ersin Yumer, Ivan Laptev, and Cordelia Schmid. Bodynet: Volumetric inference of 3d human body shapes. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 20–36, 2018.
- [72] Daniel Vlasic, Ilya Baran, Wojciech Matusik, and Jovan Popović. Articulated mesh animation from multi-view silhouettes. In *ACM Transactions on Graphics (TOG)*, volume 27, page 97. ACM, 2008.
- [73] Daniel Vlasic, Pieter Peers, Ilya Baran, Paul Debevec, Jovan Popović, Szymon Rusinkiewicz, and Wojciech Matusik. Dynamic shape capture using multi-view photometric stereo. *ACM Transactions on Graphics (TOG)*, 28(5):174, 2009.
- [74] Jiayi Wang, Franziska Mueller, Florian Bernard, Suzanne Sorli, Oleksandr Sotnychenko, Neng Qian, Miguel A. Otaduy, Dan Casas, and Christian Theobalt. RGB2Hands: Real-Time Tracking of 3D Hand Interactions from Monocular RGB Video. *ACM Transactions on Graphics (TOG)*, 39(6), 12 2020.
- [75] Shih-En Wei, Varun Ramakrishna, Takeo Kanade, and Yaser Sheikh. Convolutional Pose Machines. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- [76] X. Wei, P. Zhang, and J. Chai. Accurate realtime full-body motion capture using a single depth camera. *ACM TOG (Proc. SIGGRAPH Asia)*, 31(6):188:1–188:12, 2012.
- [77] Chenglei Wu, Carsten Stoll, Levi Valgaerts, and Christian Theobalt. On-set Performance Capture of Multiple Actors With A Stereo Camera. In *ACM Transactions on Graphics (Proceedings of SIGGRAPH Asia 2013)*, volume 32, pages 161:1–161:11, November 2013.
- [78] Donglai Xiang, Hanbyul Joo, and Yaser Sheikh. Monocular total capture: Posing face, body, and hands in the wild. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2019.
- [79] Weipeng Xu, Avishek Chatterjee, Michael Zollhöfer, Helge Rhodin, Dushyant Mehta, Hans-Peter Seidel, and Christian Theobalt. Monoperfcap: Human performance capture from monocular video. *ACM Transactions on Graphics (ToG)*, 37(2):1–15, 2018.
- [80] Genzhi Ye, Yebin Liu, Nils Hasler, Xiangyang Ji, Qionghai Dai, and Christian Theobalt. Performance capture of interacting characters with handheld kinects. In *ECCV*, volume 7573 LNCS, pages 828–841, 2012.
- [81] Mao Ye and Ruigang Yang. Real-time simultaneous pose and shape estimation for articulated objects using a single depth camera. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2345–2352, 2014.
- [82] Tao Yu, Kaiwen Guo, Feng Xu, Yuan Dong, Zhaoqi Su, Jian-hui Zhao, Jianguo Li, Qionghai Dai, and Yebin Liu. Body-fusion: Real-time capture of human motion and surface geometry using a single depth camera. In *The IEEE Interna-*

- tional Conference on Computer Vision (ICCV)*. ACM, October 2017.
- [83] Tao Yu, Zerong Zheng, Kaiwen Guo, Jianhui Zhao, Qionghai Dai, Hao Li, Gerard Pons-Moll, and Yebin Liu. Doublefusion: Real-time capture of human performances with inner body shapes from a single depth sensor. In *The IEEE International Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, June 2018.
 - [84] Tao Yu, Zerong Zheng, Yuan Zhong, Jianhui Zhao, Qionghai Dai, Gerard Pons-Moll, and Yebin Liu. Simulcap: Single-view human performance capture with cloth simulation. In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 5499–5509. IEEE, 2019.
 - [85] Zerong Zheng, Tao Yu, Yixuan Wei, Qionghai Dai, and Yebin Liu. Deephuman: 3d human reconstruction from a single image. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 7739–7749, 2019.
 - [86] Shizhe Zhou, Hongbo Fu, Ligang Liu, Daniel Cohen-Or, and Xiaoguang Han. Parametric reshaping of human bodies in images. *ACM Transactions on Graphics (TOG)*, 29(4):126, 2010.
 - [87] Xingyi Zhou, Qixing Huang, Xiao Sun, Xiangyang Xue, and Yichen Wei. Towards 3d human pose estimation in the wild: A weakly-supervised approach. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 398–407, 2017.
 - [88] Yuxiao Zhou, Marc Habermann, Weipeng Xu, Ikhsanul Habibie, Christian Theobalt, and Feng Xu. Monocular real-time hand shape and motion capture using multi-modal data. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 0–0, 2020.